

Raffaello Fornasiere

Email: fornasiere.raffaello@gmail.com LinkedIn: [raffaello-fornasiere](https://www.linkedin.com/in/raffaello-fornasiere) Website: rf-98.com GitHub: [RaffaelloFornasiere](https://github.com/RaffaelloFornasiere)

Summary

I am an AI safety researcher specializing in LLM interpretability and training methodologies. I am currently at LASR Labs, stress-testing whether current LLM interpretability methods remain reliable on models trained with different methods. I have a background in computer science and have previously worked as a lead software engineer and AI integration developer.

Professional Experience

AI Safety Researcher, LASR Labs (April 2026 - August 2026)

- My team at LASR Labs has been recently awarded a grant of £200k from Coefficient Giving to extend the project for 4 additional months, continuing under the mentorship of **Stefan Heimersheim**.
- Investigating which training variations of model organisms cause mechanistic interpretability methods to fail, to expose blind spots in current auditing benchmarks. Our paper has accepted at ICML Interpretability Workshop 2026.

Technical AI Safety course facilitator, BlueDot Impact (May 2026 - June 2026)

- Part-time (7h/week) facilitating two cohorts through BlueDot Impact's Technical AI Safety course, leading weekly discussions on AI Safety.

LASR Labs Research Fellow, LASR Labs (January 2026 - April 2026)

- Selected for a 13-week AI safety research program in London.
- Working in a team mentored by **Stefan Heimersheim** to stress test LLM interpretability methods under different training methodologies.

ARENA 7.0 Mentee, ARENA (January 2026 - February 2026)

- 5-week full-time bootcamp focused on ML engineering for technical AI safety.
- Working on transformers architecture, reinforcement learning, mechanistic interpretability, and safety evaluations.
- Awarded 1st prize at ARENA's interpretability hackathon.

AI Agent Developer, Freelance for William Saunders, Alignment Science Researcher at Anthropic (May 2025 - December 2025)

- Built a multi-agent AI assistant that maintains coherent context over months of use, avoiding the repetitive patterns typical of commercial assistants.
- Integrated tool add-ons including asynchronous browser use, Google Calendar, and speech-to-speech interfaces.
- Tech stack: FastAPI, React, Claude API, RAG, event-driven architecture.

Lead Software Engineer - Front-End, Infinitemoon (October 2022 - August 2025)

- Prototyped AI healthcare applications using LLMs, exploring feasibility for medical report automation (related to my thesis work).
- Designed, built and maintained user interfaces using Angular, ensuring seamless integration with the JHipster framework.
- Developed reusable JHipster framework extensions (e.g., an extended JHipster Criteria API, a custom Spring Data JPA repository for bidirectional relationship synchronization).
- Led the development of multiple products.
- Mentored two junior developers and promoted best practices to improve code quality.

Junior Full Stack Developer, Prodigys Group (July 2021 - September 2022)

- Developed and maintained both front-end and back-end components using Angular, SpringBoot, and Django.
- Optimized database performance through SQL query refactors, reducing execution times from hours to minutes.

Education

Master's in Computer Science and Engineering, Politecnico di Milano (Graduated July 2024)

- Thesis: Exploring the Potential of Lightweight LLMs for Medication and Timeline Extraction.
- Developed a web application integrating LLMs for medical information extraction, **tested in three Italian hospitals**.
- Published research: "Medical Information Extraction with Large Language Models" (ICNLSP-2024, ACL Anthology).

Bachelor's in Electronics Engineering, Università degli Studi di Udine (Graduated July 2020)

- Thesis: Optimization of Digital Circuit Propagation Times Using Genetic Algorithms.

Technical Skills

ML Research & Engineering:

- Model training and fine-tuning (SFT, DPO) up to 1B parameters, data pipeline construction and filtering at scale.
- Linear probes and classifiers for mechanistic interpretability, LLM-as-judge labelling pipelines.
- Transformer internals (TransformerLens, activation hooks), RL and RLHF, safety evaluations and benchmarking.
- Libraries: PyTorch, Hugging Face Transformers, TransformerLens, llama.cpp.

AI/LLM Integration:

- LLM APIs (OpenAI, Anthropic Claude), agent development, RAG, prompt engineering, local models (llama.cpp).

Languages:

- Primary: TypeScript/JavaScript, Python, Java.
- Also experienced with: C++, SQL.
- Others: CSS, HTML.

Frameworks, Libraries and Tools:

- Back-end: FastAPI, Django, Spring Boot, JHipster.
- Front-end: React, Vue.js, Angular.
- Others: Docker, Git, Tailwind CSS, Streamlit, Spring Data JPA, Postgres/PostgreSQL, SQLite.

Publications

The Model Organism Lottery: Model Organism Interpretability Strongly Depends on Training Methodology

Published at ICML Interpretability Workshop 2026

<https://arxiv.org/abs/2607.01033>

- Authors: Andrzej Szablewski, Gabriel Konar-Steenberg, Raffaello Fornasiere, Nikita Menon, Stefan Heimersheim

Medical Information Extraction with Large Language Models

Published at ICNLP-2024, ACL Anthology

<https://aclanthology.org/2024.icnlp-1.47/>

- Authors: **Raffaello Fornasiere**, Nicolò Brunello, Vincenzo Scotti, Mark Carman.

Academic Projects

Medical Image-Caption Matching with Transformers

- Reimplemented CLIP model from scratch for medical radiology applications using ResNet vision encoder and custom transformer text encoder.
- Trained and fine-tuned both models using contrastive learning, pre-trained BERT integration, caption preprocessing, and medical concept augmentation techniques.
- Developed evaluation methods for assessing image-caption alignment in medical domain.

Online Learning for Dynamic Pricing in E-commerce

- Exploration work implementing multi-armed bandit algorithms (GP-TS, GP-UCB) with social influence modeling to optimize advertising budget allocation across product campaigns.

Deep Learning Projects with CNNs

- **Species Classification:** Implemented transfer learning with VGG16/19 and EfficientNet models, applying data augmentation techniques (CutMix, MixUp) and ensemble methods to achieve 86.91% accuracy on image classification tasks.
- **Time Series Classification:** Developed CNN-LSTM hybrid architectures for sequential data analysis, experimenting with preprocessing techniques including normalization and noise augmentation to reach 70% classification accuracy.

Other Experiences

STAIR Mentor, St Andrews Impact Research (Feb 2026 - March 2026)

- Mentoring 2 early-career students beginning AI safety research.

ML4Good AI Safety Bootcamp

- Completed EPFL's iteration of the ML4Good AI Safety and ML Bootcamp (February 2023).

Community Involvement

Volunteer developer for hometown community

- Built website to publish local competition status with real-time updates through Telegram bot.
- Developed web application to emulate Italian TV game show, managing multiplayer gameplay across mobile devices.